

Classification of regulatory sequences using machine learning techniques

Boris Shilov, Wim Thiels, Lucas Coppens, Zixuan Xie

Introduction

Melanoma is an aggressive cancer with a high level of therapy-resistance due to its high degree of heterogeneity and plasticity. Despite the great efforts put into the field of melanoma treatment, resulting in great advances, many challenges remain. The major challenge in fighting melanoma relapse continues to be the tumor heterogeneity, as melanoma comprises a wide variety of phenotypically distinct subpopulations of cancer cells. To be able to address this issue therapeutically, the underlying mechanisms of heterogeneity need to be characterized. In this paper, a distinction between two major types of melanoma cells is made: invasive and proliferative. (Verfaillie et al. 2015; Shannan et al. 2015)

Transcriptional reprogramming of melanoma cells in proliferative state into melanoma cells with invasive characteristics is a critical event at the origin of metastatic spreading of melanoma. Invasive cells have acquired the ability to migrate to other tissues, enter the bloodstream and therefore lie at the basis of the metastatic spreading of cancer in the body. While the transitional mechanisms from proliferative to invasive cancer cell are yet to be characterized more extensively, it is sure that one event lies at the basis of this transition: the transcriptional reprogramming of the cell. Studying the involved genes and regulatory elements using various bioinformatics approaches is therefore a hot topic in the area of melanoma research. Decoding the regulatory landscape could result into the ability to push melanoma cells towards a different cell state, which would be an interesting target from a therapeutic point of view. (Verfaillie et al. 2015)

Transcriptomic, open chromatin and histone modification maps of melanoma cultures were constructed, revealing thousands of active cis-regulatory regions, both for proliferative and invasive cells. (Verfaillie et al. 2015) It should be possible to discern which cis-regulatory regions are useful for the classification of cell states in melanoma samples if such states are truly distinct in terms of regulatory landscape. The aim of the project was the construction of classifiers predicting whether a regulatory region would be active in proliferative or invasive cell states. Another useful insight that would be gained from these classifiers, is which regulatory regions are the most significant for distinguishing between cell states, thus giving informa-

tion about the underlying mechanisms and the critical genes and regulatory elements involved in cancer cell state transitions.

Two distinct machine learning techniques were used for the creation of such classifiers, namely the random forests ensemble method and deep learning, with the use of convolutional neural networks. Random forest (RF) is a nonparametric tree-based method that builds an ensemble model from random subsets of features. (Nguyen et al. 2015) RF has shown excellent performance for classification problems, it works well when the number of features is much larger than the number of samples. However, with randomizing mechanism in feature selection, RF could give poor accuracy when applied to high dimensional data. This is mainly caused by the subspace of features randomly sampled from hundreds of features to split a node of the tree is often dominated by uninformative features (or noise) when growing a tree from the data. In this manner the tree grown from such subspace of features will result in a low accuracy RF model. (Nguyen et al. 2015) Aiming at improving the performance of the RF model, we propose a feature selection algorithm before constructing the RF model, called Boruta. Boruta is an all-relevant feature selection method. It tries to capture all the important, interesting features that might be relatively important to the outcome variable. (Kursa et al. 2010) In addition to the RF method, we also build convolutional neural networks parallelly. A neural network is able to extract features from raw DNA sequences. No prior manual feature selection is necessary. By converting the sequence into a 1-hot vector, it becomes an image-like input which can be scanned with a number of kernels from the first convolutional layer. This information is then passed on to the rest of the network. During training, the network will adjust its weights in order to optimize the classification task. In the process, the kernel weights of the first convolutional layer will start to resemble the various DNA motifs which are then extracted and compared to a motif database (JASPAR). (Min et al. 2017) Both models were trained on the same training set, which comprises the dataset of active cis-regulatory regions, mentioned hereabove. In this paper, both methods are described and their results are evaluated and compared.

Methods

Random forests

Unlike neural networks, random forests do not have feature learning integrated into the training process. Thus, they require the input data to have labelled features. Before the application of RFs to our data was possible, we needed a way to discover the features. This limitation also precluded us from training our model on the sequence data.

We therefore had to engineer a pipeline to process the sequence data into a usable form, which took a considerable investment of time. Our reasoning on this front was as follows: the sequences themselves are not the principal target of our learning strategy, but rather whatever recurring motifs may be present.

To this end, we first used the motif discovery utility Hypergeometric Optimization of Motif EnRichment (HOMER), which provides a Perl module `findMotifs.pl` that allows for the discovery of enriched regulatory motifs in one set of sequences as compared to another set (the background) (Heinz et al.). This enabled us to find these motifs in the sequences, and crucially, because we want to compare and contrast the Invasive and Proliferative cell states, allowed us to run these two sets of sequences as backgrounds against each other, identifying the most differential motifs and increasing our confidence in the next steps. The results of these two runs can be found in `HomerOutput`. HOMER analysis both found known motifs and discovered some motifs *de novo*. Every motif is defined in terms of a position weight matrix (PWM).

Having identified the ensemble of differentially enriched motifs in these two cell states compared to each other, we used Cluster-Buster in order to score the Invasive and Proliferative sequence sets for each motif, independently (Frith, Li, and Weng –). This resulted in a log-likelihood score for every motif against every sequence in the sequence sets. The functions for this step of the analysis can be found in `motif_processing_main.sh`. We proceeded to assemble these into a feature matrix, with every row being the genomic region and every column a particular motif, hereafter referred to as features, with the log-likelihood scores as values in this matrix. We further inserted a `_label` column which represented our binary outcome variable, with 0 for Invasive state and 1 for Proliferative state. This step of the analysis can be found in `feature_matrices.py`, utilising the `feature_matrix_special` method of the `cbust_result` class we created in order to handle the output of Cluster-Buster. We thus produced a complete data

matrix for our subsequent random forest classifier training and exploration.

Random forests

Random forest (RF) is a nonparametric tree-based method that builds an ensemble model from random subsets of features. (Nguyen et al. 2015) RF has shown excellent performance for classification problems, it works well when the number of features is much larger than the number of samples. However, with its typical randomizing mechanism in feature selection, RF models risk having a poor performance when applied to high dimensional data. The main reason for this is that the RF ensemble method uses subset of features randomly sampled from hundreds of features in each iteration. If such a subset is poorly chosen, the choice of nodes for each decision tree is often dominated by uninformative features. This way, trees grown from such a subspace of uninformative features will result in a RF model yielding poor accuracy. (Nguyen et al. 2015) Aiming at improving the performance of the RF model, a feature selection algorithm filtering out the uninformative features before constructing the RF model was proposed. Boruta is an all-relevant feature selection method. It tries to capture all the relevant features that might be relatively important to the outcome variable. (Kursa et al. 2010)

Construction of the feature matrix

As random forests do not have feature learning integrated into the training process. Thus, they require the input data to have labelled features. Before the application of RFs to our data was possible, a way to discover motifs was required, which would then be the features used in training the RF model.

As the raw sequence data was not suitable for training the random forests model, a pipeline to process the sequence data into a usable form, i.e. a feature matrix, was engineered, taking a considerable amount of time. However this investment was justified by the fact that the sequences themselves were not the principal target, but rather whatever recurring motifs would be present.

Motif discovery and scoring

To discover motifs in the input fasta sequences, the motif discovery utility Hypergeometric Optimization of Motif EnRichment (HOMER) was used, providing a Perl module `findMotifs.pl` that allows for the discovery of enriched regulatory motifs in one set of sequences as compared to another set (the back-

ground) (Heinz et al.). As the goal was to compare and contrast the invasive and proliferative cell states, these two sets of sequences were ran as backgrounds against each other, identifying the most differential motifs. The results of these two runs can be found in `HomerOutput`. HOMER analysis served two purposes: it found known motifs and discovered some *de novo* motifs. The motifs returned by HOMER were formatted as position weight matrices (PWM), representing the relative frequencies of each nucleic acid on each position in the motif.

Having identified the ensemble of differentially enriched motifs in these two cell states compared to each other, the Cluster-Buster software was used in order to score the invasive and proliferative sequence sets for each motif, independently (Frith, Li, and Weng –). This resulted in a log-likelihood score for every motif against every sequence in the sequence sets. The functions for this step of the analysis can be found in `motif_processing_main.sh`. The set of sequences scored for every motif was then assembled into a feature matrix. In such a feature matrix, every row represents a genomic region and every column represents a particular motif, hereafter referred to as features, with the log-likelihood scores as values in this matrix. Furtherly, a `_label` column representing the binary outcome variable was inserted, with value 0 representing the invasive cell state and value 1 representing the proliferative cell state. This step of the analysis can be found in `feature_matrices.py`, utilising the `feature_matrix_special` method of the `cbust_result` class, created in order to handle the output of Cluster-Buster. Thus, a complete data matrix was produced, formatted correctly to be used as training data for a random forests classifier.

Feature Selection

MAST AND BORUTA

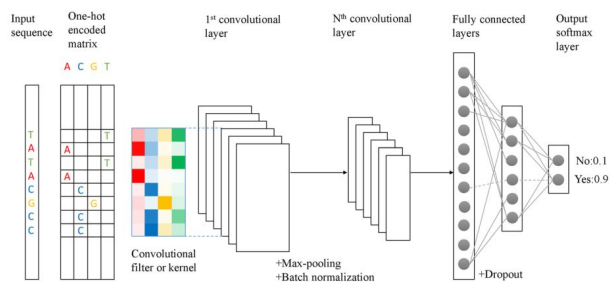
Construction of the model

TEXT ABOUT THE SCIKIT LEARN PYTHON PACKAGE

Deep learning

2 architectures

A schematic outline of how a neural network can be used in motif detection is given in figure ??.



In our case, 2 model architectures were compared (see fig??). Model 1 uses 3 CNN's, Model 2 uses only 1 convolutional layer, and incorporates a recurrent layer (LSTM's). A recurrent neural layer can pick up on sequential information and are therefore better suited to capture the context of the motifs (motif syntax).

Model 1			Model 2		
Layer (type)	Output Shape	Param #	Layer (type)	Output Shape	Param #
conv1d_1 (Conv1D)	(None, 797, 30)	2310	conv1d_1 (Conv1D)	(None, 790, 32)	3360
max_pooling1d_1 (MaxPooling1D)	(None, 199, 30)	0	dropout_1 (Dropout)	(None, 790, 32)	0
dropout_3 (Dropout)	(None, 199, 30)	0	time_distributed_lstm (TimeDistributedLSTM)	(None, 790, 32)	1056
conv1d_2 (Conv1D)	(None, 189, 20)	6420	max_pooling1d_1 (MaxPooling1D)	(None, 60, 32)	0
max_pooling1d_2 (MaxPooling1D)	(None, 47, 20)	0	bidirectional_1 (Bidirectional)	(None, 60, 64)	16640
dropout_2 (Dropout)	(None, 47, 20)	0	dropout_2 (Dropout)	(None, 60, 64)	0
conv1d_3 (Conv1D)	(None, 41, 20)	2020	flatten_1 (Flatten)	(None, 3040)	0
max_pooling1d_3 (MaxPooling1D)	(None, 10, 20)	0	dense_2 (Dense)	(None, 64)	245824
dropout_3 (Dropout)	(None, 10, 20)	0	dropout_3 (Dropout)	(None, 64)	0
flatten_1 (Flatten)	(None, 200)	0	dense_3 (Dense)	(None, 2)	130
dense_1 (Dense)	(None, 925)	185925			
activation_3 (Activation)	(None, 925)	0			
dropout_4 (Dropout)	(None, 925)	0			
dense_2 (Dense)	(None, 2)	1852			
Total params: 199,527			Total params: 267,010		

Both models were build using the python Keras package on a Tensorflow backend.

Data augmentation

Even though a neural network can work with raw DNA sequences as input, there are still two issues that need to be overcome. First, the input to the convolutional layer needs to be of fixed length, and second, a deep neural network needs a vast amount of training data. To overcome these 2 issues, a moving window approach as in (Min et al. 2017) (with stride = 20) is used (fig ??).

training/validation

Starting from 20122 enhancer sequences, we end up with 228051 sequences using the data augmentation approach. The sequences are randomly split in a 70/20/10 fashion (train/test/val). The loss function was class weighted to counter the fact that the I-label is overrepresented in the input data. Both the validation and test data were also balanced.

motif extraction

After training, the kernels of the first convolutional layer, can be screened for motifs. To go from the kernel

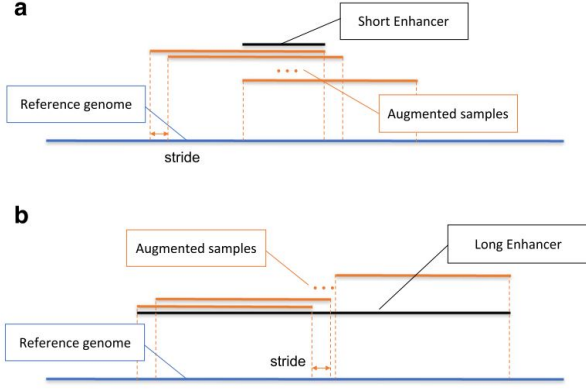


Figure 1: Diagram of data augmentation. By using a sliding window approach, each enhancer sequence produces multiple sequences of equal length that can be used as input for the first convolutional layer. Taken from Min, Xu et al (2017)




Figure 2: Training metrics for both models

	accuracy	AUROC	AUPR	Epoch time
Model 1	0.69	0.75	0.66	10min
Model 2	0.60	0.64	0.61	15min

Figure 3: Comparing classification performance of both models




Figure 4: ROC and Precision Recall curves comparing both models

weights to a proper PSW matrix for a motif, each of the kernels are scored against each of the original DNA sequences. This score represents the best match of a motif when slid across a particular sequence. For each motif, the top 100 scoring sequences are stored, and these 100 sequences are then summarized into 1 PSWM. These PSWM's are then queried against the JASPAR motif database using TOMTOM [2]. Only motif matches with a threshold of $E < 0.01$ are considered significant.

Results and discussion

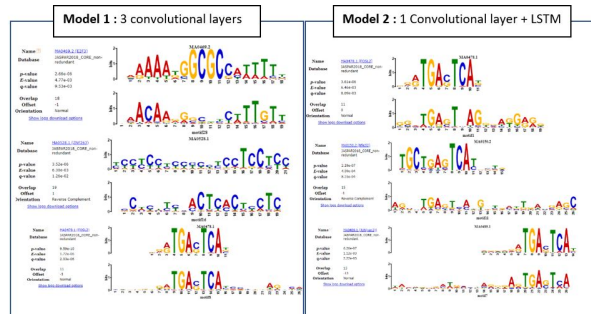
Deep learning

Comparing classification performance

Both models are tested on a balanced test set of 45333 sequences (20% of the total). From Table1 and Figure?? it is clear that model 1 outperforms model 2. The model using multiple convolutional layers achieves an accuracy that is 9% higher then the model using the recurrent layer.

motifs detected

The motifs detected by both models are shown in fig ??



References

Frith, Martin C, Michael C Li, and Zhiping Weng. —. “Cluster-Buster: Finding Dense Clusters of Motifs in Dna Sequences.” *Nucleic Acids Res* 31 (13). Bioinformatics Program, Boston University, 44 Cummington Street, Boston, MA 02215, USA.: 3666–8.

Heinz, Sven, Christopher Benner, Nathanael Spann, Eric Bertolino, Yin C Lin, Peter Laslo, Jason X Cheng, Cornelis Murre, Harinder Singh, and Christopher K Glass “Simple Combinations of Lineage-Determining Transcription Factors Prime Cis-Regulatory Elements Required for Macrophage and B Cell Identities.” *Mol Cell* 38 (4). Department of Cellular; Molecular Medicine, University of California, San Diego, La Jolla, CA 92093, USA.: 576–89. doi:10.1016/j.molcel.2010.05.004.

Kursa et al. 2010. “Feature Selection with the Boruta Package.” *Journal of Statistical Software*.

Min et al. 2017. “Predicting Enhancers with Deep Convolutional Neural Networks.” *BMC Bioinformatics*.

Nguyen et al. 2015. “Unbiased Feature Selection in Learning Random Forests for High-Dimensional Data.” *The Scientific World Journal*.

Shannan et al. 2015. “Heterogeneity in Melanoma.” *Cancer Treatment and Research*.

Verfaillie et al. 2015. “Decoding the Regulatory Landscape of Melanoma Reveals Teads as Regulators of the Invasive Cell State.” *Nature Communications*.